

CS127 Quiz #2

Date: November 13th, 2019 at 3:00 pm

Your BannerID

BannerID: _____

Write your solutions on this piece of paper and hand it in.

Question #1 General Topics

30 points

- A. Unlike ISAM, _____, and _____ dynamically adjust to inserts and deletes.
- a. clustered indexes, non-clustered indexes
 - b. dense indexes, sparse indexes
 - c. sequential indexes, static hash indexes
 - d. B+-Trees, extendible hashing indexes
 - e. slotted page structures, multitable clusterings
- B. _____ allow for better compression of data in memory than _____.
- a. B-Trees, sequential indexes
 - b. multitable clusterings, sequential indexes
 - c. queries that involve joins, queries that do not involve joins
 - d. B+-Trees, B-Trees
 - e. columnstores, rowstores
- C. The buffer manager reads blocks into the buffer. When a transaction accesses a block, the buffer manager _____ the block in memory to prevent the block from being written back to disc.
- a. sorts
 - b. deletes
 - c. pins
 - d. flushes
 - e. encrypts

- D. _____ enhance performance of queries that involve _____, and decrease performance of queries that involve _____.
- a. multitable clusterings, joins, single relations
 - b. free lists, deletions, insertions
 - c. hash indexes, ranges, returning a single tuple
 - d. B+-Trees, returning a single tuple, ranges
 - e. dense indexes, joins, single relations
- E. Given a candidate key attribute, _____ on the attribute may contain exactly one or two copies of the same search key, while _____ on the attribute may contain only one copy of the same search key.
- a. dense indexes, sparse indexes
 - b. non-clustered indexes, clustered indexes
 - c. extendible hashing indexes, linear hashing indexes
 - d. B-Trees, dense indexes
 - e. B+-Trees, B-Trees
- F. _____ are well-suited for fixed length tuples, and _____ are well-suited for variable-length tuples.
- a. sequential storage structures, slotted page storage structures
 - b. hash indexes, B+-Tree indexes
 - c. dense indexes, sparse indexes
 - d. primary indexes, secondary indexes
 - e. merge joins, indexed nested loop joins

Question #2 Dense and Sparse Indices

10 points

A sequential file (sorted by the key field) consists of 10,000 records. Blocks are 1000 bytes long; there is no need for a block header. Records are 100 bytes long, of which 12 bytes are the key field. Pointers take 8 bytes.

- A. Assuming there is an infinite amount of main memory, the number of blocks required for a sequential dense index on a primary key in this file is (Hint: Don't think B-tree)
- a 80
 - b 120
 - c 200
 - d 800
 - e 1200
- B. The number of blocks required for a block-oriented, sparse index on this file is
- a 8
 - b 12
 - c 20
 - d 80
 - e 120

Question #3 B+-Trees

10 points

- A. A B+-tree of order 4 and of height 3 will have a maximum of _____ leaves.
- a. 64
 - b. 128
 - c. 186
 - d. 256
 - e. 512
- B. Consider a B+-Tree of order 3 (each node has 3 children, and 2 keys), the tree here shows the **keys** of the tree:




Figure 1: B+-Tree Insertion Problem

What is the height of the tree after inserting 3 then 7 (in that order)?

- a. 0
- b. 1
- c. 2
- d. 3
- e. 4

Question #4 Extendable Hashing

15 points

With x as the employee name, hash $h(x)$ returns the following 10-bit values:

| x | $h(x)$ |
|-------|------------|
| Pete | 1001111010 |
| Mary | 0100000000 |
| Jane | 1100011110 |
| Bill | 0110000000 |
| John | 0001101001 |
| Vince | 1101110101 |
| Karen | 0000110111 |

Using the above values, and a bucket size of 2, we draw a picture of the state of an extendable hash table after records with the following keys are inserted in the given order:

Pete, Karen, Mary, Vince, Jane, John, Bill




Figure 2: Extendable Hashing problem

- A. Does this picture change if we insert the keys in the opposite order?
- B. Draw an altered picture to show what happens when we add a record for Sol to the extendable hash table in Figure 2 if $h(\text{Sol}) = 0001010001$.

Question #5 Linear Hashing

15 points

Your linear hash table uses a hash function $h(x) = x$. Each bucket in the table has a size = 2. The table starts with 2 buckets. After hashing 5, 4, 6, and 3, it looks like this. Note the split pointer is pointing at bucket 0.




Figure 3: Linear Hashing problem

- A. Draw the result after inserting 2. Your drawing should show the final location of the split pointer after the operation is completed.
- B. Draw the result after inserting 10. Your drawing should show the final location of the split pointer after the operation is completed.

Question #6 Query Processing

20 points

Consider

- Relations $R(A, B)$ and $S(A, C)$ with the b_r and b_s number of blocks, and n_r and n_s numbers of records, respectively.
 - Neither R nor S fit in memory
- A. Block nested loop join for relations R and S has two for loops: an inner and outer. If main memory could fit M blocks, and there are no indexes, how many iterations does the outer loop go through?
- B. R has a clustered index on A (implying that the relation's tuples are sorted on A). If we wanted to use the index to satisfy the query `SELECT * from R where A > 10`, how many times do we need to traverse the index?
- C. If S also has a clustered index on A , what is the lowest-cost way of joining R and S on A ?